

Notes: `df.ix[:]` : returns sliced rows based on passed indexes

Predicting US Stocks Returns Direction (UP/DOWN)

Implementing Anchor bias theory

Problem Identification :

Investors are anchored to the long term moving averages. The long term moving average is defined by the 252 moving average, and the short term is defined by the 21-Day moving average. The distance between the two moving averages is the moving average distance ($MAD = 21\text{-DAY MA} / 252\text{-DAY MA}$). When the $MAD > 1$, the distance is called a positive spread and when $MAD < 1$, the distance is called a negative spread.

The anchor bias theory, published in a research paper by Avramov, Kaplanski, Subrahmanyam(2018), states that when MAD spread is positive positive announcement (sentiment) drive the price of the stocks to go up more than negative sentiment drive the price to go down. However, when MAD spread is negative, negative sentiment drives price to go down more than positive sentiment drives re price to go up. Noting that the larger/smaller the MAD, in both cases, the more effective is the strategy

The model proposed is to predict US stocks returns (+/-) based on several features but mainly on a BUY or SELL signal. The engineered feature, named trading signal is the main feature which is processed by the constructed pipeline. The BUY signal is constructed by getting positive sentiment from 2 databases (Sentdex and stocktwits), a 7 days previous sentiment score and a positive MAD greater than 1.2. The SELL signal is set based on negative sentiment scores from 2 databases, also a 7 day previous negative score and a negative MAD less than 0.8.

The stated signals are passed to the pipeline to pass through more than 8000 US stocks and filter out each day, the stocks that passed the criteria. Several screens were passed to the timeline to insure no stock has a null sentiment score (in any of the two databases) or a zero return (which was actually found). Several other features were passed to the pipeline to output a dataframe of the filtered stocks. After doing the necessary transformations, the data is based to two machine learning algorithms.

Data Gathering using a Pipeline:

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
sns.set_style('white')
```

```

In [2]: # Import Pipeline class and datasets
from quantopian.pipeline import Pipeline
from quantopian.pipeline.data import USEquityPricing
from quantopian.pipeline.data.psychsignal import stocktwits
# from quantopian.interactive.data.sentdex import sentiment
from quantopian.pipeline.data.sentdex import sentiment
from quantopian.pipeline.factors import CustomFactor, MarketCap, Latest
from quantopian.pipeline.classifiers.fundamentals import SuperSector
# Import built-in moving average calculation
from quantopian.pipeline.factors import SimpleMovingAverage, DailyReturns, Returns

# Import built-in trading universe
from quantopian.pipeline.experimental import QTradableStocksUS

# Define our own custom factor class
class SentimentSevenDaysAgo(CustomFactor):
    inputs = [sentiment.sentiment_signal]
    window_length=7

    def compute(self, today, assets, out, sentiment):
        out[:] = sentiment[0] #When I specified a window_length of 7 it gave back last 7 scores thus call the sentiment Boundcolumn from the top index

def make_pipeline():
    # Create a reference to our trading universe
    base_universe = QTradableStocksUS()

    # Get latest closing price
    close_price = USEquityPricing.close.latest

    daily_returns = DailyReturns(
        inputs = [USEquityPricing.close])

    returns_3 = Returns(
        inputs = [USEquityPricing.close],
        window_length = 2)

    mean_close_21 = SimpleMovingAverage(
        inputs = [USEquityPricing.close],
        window_length= 21)

    mean_close_252 = SimpleMovingAverage(
        inputs = [USEquityPricing.close],
        window_length= 252)

    MAD = mean_close_21 / mean_close_252

    sentdex_score = sentiment.sentiment_signal.latest

    sentdex_lag = SentimentSevenDaysAgo()

    marketCap = MarketCap()

```

```

stock_class = SuperSector()

not_zero_returns = (daily_returns != 0) & (returns_3 !=0)
# Calculate 3 day average of bull_minus_bear scores
sentiment_score = SimpleMovingAverage(
    inputs=[stocktwits.bull_minus_bear],
    window_length=2,
)
# Create filter for positive/negative spread moving averages
# Create filter for positive sentiment scores
# Create filter for 7 days lag sentiment score
# assets based on their sentiment scores
positive_MAD = MAD > 1.2
negative_MAD = MAD < 0.8

positive_sentiment_lag = sentdex_lag > 3
negative_sentiment_lag = sentdex_lag < -1

positive_sentiment = sentdex_score > 2
negative_sentiment = sentdex_score < -1

positive_twits = sentiment_score > 0
negative_twits = sentiment_score < 0
# Long = sentdex_lag.top(10, mask = positive_MAD)
# short = sentdex_lag.bottom(10, mask = negative_MAD)

Long = (positive_MAD & positive_sentiment_lag & positive_twits & positive_
sentiment)
short = (negative_MAD & negative_sentiment_lag & negative_twits )

tradeable_equities = (Long | short)
# Return Pipeline containing all below columns and
# sentiment_score that has our trading universe as screen
return Pipeline(
    columns={
        'close_price': close_price,
        "Sentdex": sentdex_score,
        "Sentdex_lag": sentdex_lag,
        'sentiment_score': sentiment_score.zscore(), #apply zscore to norm
alize
        "MAD": MAD,
        "BUY": Long,
        "SHORT": short,
        "return": daily_returns,
        "Returns": returns_3,
        "Market Capital.": marketCap,
        "Stock Classfiaction": stock_class
    },
    screen=(base_universe
        & tradeable_equities& sentdex_lag.notnull() & sentiment_score.notnull
        () & not_zero_returns)
    )

```

```
In [3]: # Import run_pipeline method
        from quantopian.research import run_pipeline

        # Specify a time range to evaluate
        period_start = '2013-01-01 07:12:03.6'
        period_end = '2019-01-01 07:12:03.6 '

        # Execute pipeline created by make_pipeline
        # between start_date and end_date
        pipeline_output = run_pipeline(
            make_pipeline(),
            start_date=period_start,
            end_date=period_end
        )
        # pipeline_output.add(sentiment_free.sentiment_signal, 'sentiment_signal')

        # Display last 10 rows
        pipeline_output.tail(20)
        # print('Number of securities that passed the filter: %d' % len(pipeline_output))
```

Pipeline Execution Time: 7 Minutes, 12.20 Seconds

Out[3]:

		BUY	MAD	Market Capital.	Returns	SHORT	Sentdex	Sentdex
2018-12-24 00:00:00+00:00	Equity(337 [AMAT])	False	0.743352	2.905535e+10	-0.019417	True	-3.0	
	Equity(4705 [MKC])	True	1.257794	1.818684e+10	-0.007614	False	5.0	
	Equity(7447 [TIF])	False	0.790511	9.237367e+09	-0.029694	True	-3.0	
2018-12-26 00:00:00+00:00	Equity(4705 [MKC])	True	1.252156	1.767615e+10	-0.029099	False	5.0	
	Equity(7447 [TIF])	False	0.779504	9.045988e+09	-0.021633	True	-3.0	
	Equity(337 [AMAT])	False	0.735128	2.937169e+10	0.057971	True	-1.0	
2018-12-27 00:00:00+00:00	Equity(3149 [GE])	False	0.568228	6.427907e+10	0.066378	True	-3.0	
	Equity(5029 [MRK])	True	1.207782	1.924279e+11	0.038607	False	6.0	
	Equity(13635 [DO])	False	0.653096	1.330368e+09	0.056831	True	-3.0	
2018-12-28 00:00:00+00:00	Equity(39778 [QEP])	False	0.698206	1.337756e+09	0.108930	True	2.0	
	Equity(5029 [MRK])	True	1.206543	1.960164e+11	0.018654	False	6.0	
	Equity(9883 [ATVI])	False	0.690358	3.589391e+10	0.013569	True	-3.0	
2018-12-31 00:00:00+00:00	Equity(39778 [QEP])	False	0.685779	1.361433e+09	0.015929	True	6.0	
	Equity(3149 [GE])	False	0.570777	6.532284e+10	0.034388	True	-3.0	
	Equity(5029 [MRK])	True	1.204639	1.959904e+11	0.000133	False	6.0	
2019-01-02 00:00:00+00:00	Equity(39778 [QEP])	False	0.673540	1.325917e+09	-0.025261	True	6.0	
	Equity(3149 [GE])	False	0.571820	6.584473e+10	0.005984	True	-3.0	
	Equity(4705 [MKC])	True	1.234640	1.832881e+10	0.002015	False	5.0	
	Equity(5029 [MRK])	True	1.202823	1.986948e+11	0.014727	False	6.0	
	Equity(13635 [DO])	False	0.634035	1.297384e+09	-0.031828	True	-3.0	


In [4]: `pipeline_output.describe()`

Out[4]:

	MAD	Market Capital.	Returns	Sentdex	Sentdex_lag	Stock Classfiacion	close_l
count	7721.000000	7.721000e+03	7721.000000	7721.000000	7721.000000	7721.000000	7721.00
mean	1.169052	2.067004e+10	0.000524	3.778267	3.717394	2.036006	62.20
std	0.280565	3.440918e+10	0.026242	3.194239	3.488942	0.883938	95.30
min	0.212168	2.632199e+08	-0.332248	-3.000000	-3.000000	1.000000	0.91
25%	1.204733	5.088438e+09	-0.010000	4.000000	4.000000	1.000000	22.62
50%	1.239767	1.209133e+10	0.001032	5.000000	6.000000	2.000000	41.93
75%	1.302449	2.334478e+10	0.011609	6.000000	6.000000	3.000000	70.46
max	2.138074	7.399760e+11	0.528736	6.000000	6.000000	3.000000	1370.43

Pipeline Schema to Fetch US Tradeable Equities

In [5]: `# pipeline_output['return'] = pd.get_dummies(pipeline_output['return'],drop_first=True)`
`# pipeline_output.head(30)`
`make_pipeline().show_graph(format='jpeg') #or png`



Feature Engineering

In []:

In [4]: `# return_encoding = pd.get_dummies(pipeline_output['return'],drop_first=True)`
`# pipeline_output['return'] = return_encoding`
`# pipeline_output['Returns'] = pipeline_output["Returns"].apply(np.sign)`

`pipeline_output['return'] = pipeline_output["return"].apply(np.sign)`
`pipeline_output['Market Capital.'] = pipeline_output["Market Capital."].apply(np.log)`
`# Applied zscore to Market Capitalization but got an outlier which was Apple s tock`

In [5]: `pipeline_output.head(10)`

Out[5]:

		BUY	MAD	Market Capital.	Returns	SHORT	Sentdex	Sentdex_l
2013-01-03 00:00:00+00:00	Equity(754 [BBY])	False	0.630091	22.111144	-0.004219	True	-3.0	-3
2013-01-04 00:00:00+00:00	Equity(754 [BBY])	False	0.629542	22.111144	0.013559	True	-1.0	-3
	Equity(3645 [HOV])	True	1.890448	20.662923	0.010130	False	3.0	6
	Equity(754 [BBY])	False	0.631212	22.111144	0.011706	True	-1.0	-3
2013-01-07 00:00:00+00:00	Equity(3212 [GILD])	True	1.307798	24.744883	0.010538	False	6.0	6
	Equity(32902 [FSLR])	True	1.287577	21.712406	-0.024695	False	6.0	4
2013-01-08 00:00:00+00:00	Equity(3212 [GILD])	True	1.306559	24.744883	0.014916	False	6.0	4
	Equity(3645 [HOV])	True	1.918963	20.662923	-0.043353	False	3.0	6
2013-01-09 00:00:00+00:00	Equity(3212 [GILD])	True	1.306288	24.744883	0.005202	False	6.0	4
2013-01-10 00:00:00+00:00	Equity(14848 [AABA])	True	1.218424	23.823049	-0.018293	False	3.0	6

```
In [6]: pipeline_output["Trading Signal"] = pd.get_dummies(pipeline_output['BUY'],drop
_first=True)

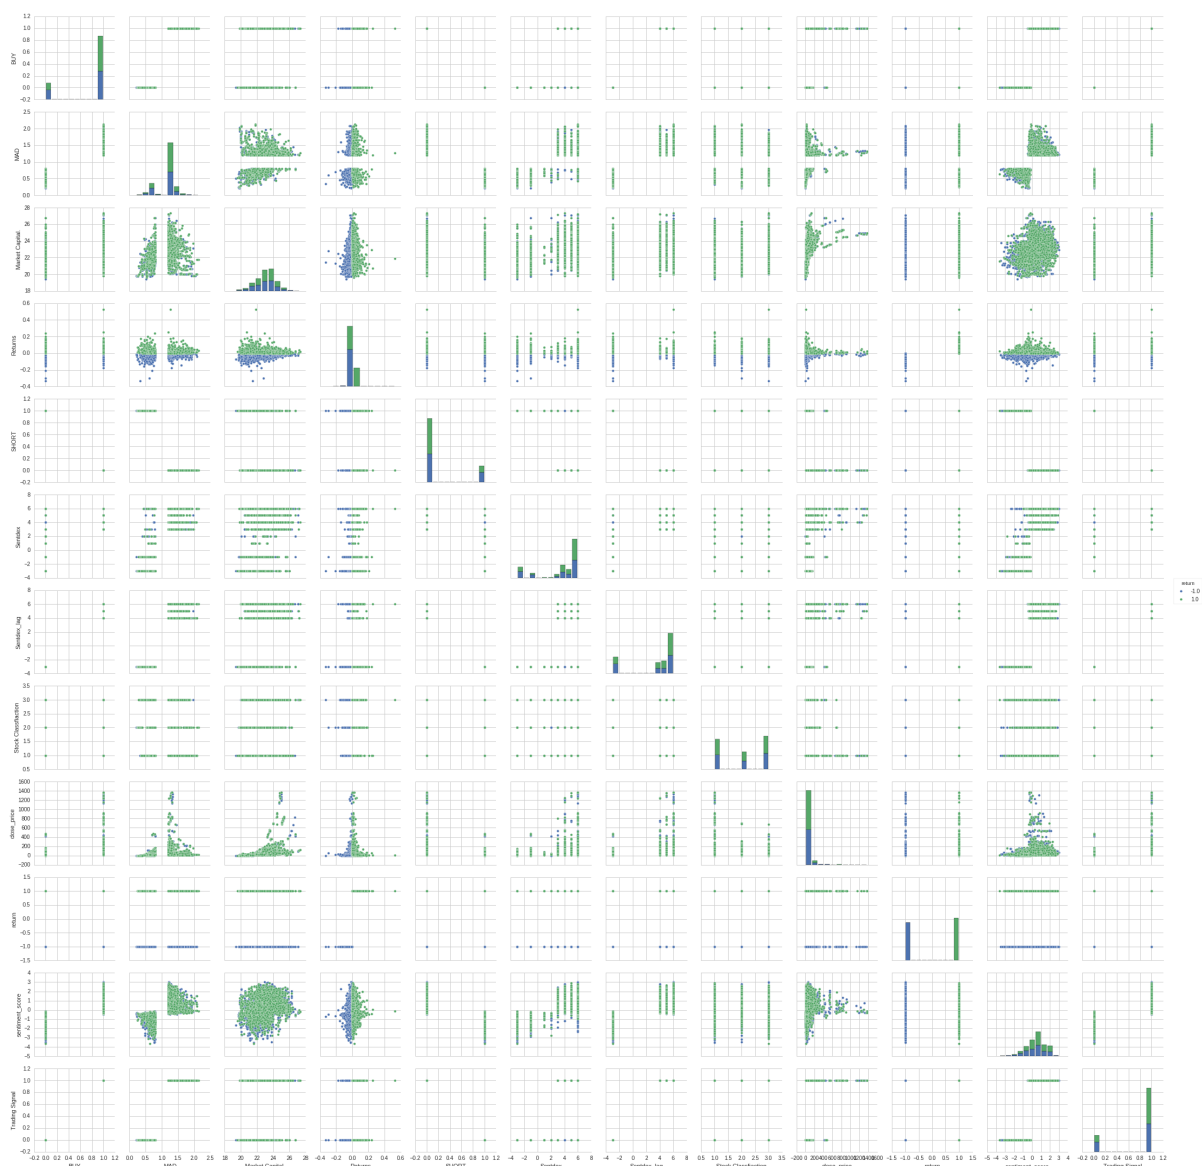
pipeline_output.tail(20)
# if BUY "FALSE" --> 0 | if BUY "TRUE" --> 1
```

Out[6]:

		BUY	MAD	Market Capital.	Returns	SHORT	Sentdex	Sentdex_k
2018-12-24 00:00:00+00:00	Equity(337 [AMAT])	False	0.743352	24.092468	-0.019417	True	-3.0	-3
	Equity(4705 [MKC])	True	1.257794	23.623964	-0.007614	False	5.0	5
	Equity(7447 [TIF])	False	0.790511	22.946523	-0.029694	True	-3.0	-3
2018-12-26 00:00:00+00:00	Equity(4705 [MKC])	True	1.252156	23.595482	-0.029099	False	5.0	5
	Equity(7447 [TIF])	False	0.779504	22.925587	-0.021633	True	-3.0	-3
	Equity(337 [AMAT])	False	0.735128	24.103297	0.057971	True	-1.0	-3
2018-12-27 00:00:00+00:00	Equity(3149 [GE])	False	0.568228	24.886500	0.066378	True	-3.0	-3
	Equity(5029 [MRK])	True	1.207782	25.982987	0.038607	False	6.0	6
	Equity(13635 [DO])	False	0.653096	21.008721	0.056831	True	-3.0	-3
2018-12-28 00:00:00+00:00	Equity(39778 [QEP])	False	0.698206	21.014259	0.108930	True	2.0	-3
	Equity(5029 [MRK])	True	1.206543	26.001464	0.018654	False	6.0	6
	Equity(9883 [ATVI])	False	0.690358	24.303833	0.013569	True	-3.0	-3
2018-12-31 00:00:00+00:00	Equity(39778 [QEP])	False	0.685779	21.031804	0.015929	True	6.0	-3
	Equity(3149 [GE])	False	0.570777	24.902608	0.034388	True	-3.0	-3
	Equity(5029 [MRK])	True	1.204639	26.001331	0.000133	False	6.0	6
2019-01-02 00:00:00+00:00	Equity(39778 [QEP])	False	0.673540	21.005370	-0.025261	True	6.0	-3
	Equity(3149 [GE])	False	0.571820	24.910565	0.005984	True	-3.0	-3
	Equity(4705 [MKC])	True	1.234640	23.631740	0.002015	False	5.0	5
	Equity(5029 [MRK])	True	1.202823	26.015036	0.014727	False	6.0	6
	Equity(13635 [DO])	False	0.634035	20.983615	-0.031828	True	-3.0	-3


```
In [14]: sns.pairplot(pipeline_output, hue='return')
```

```
Out[14]: <seaborn.axisgrid.PairGrid at 0x7f1772eb93d0>
```



Data Analysis and Insights Generation

```
In [9]: n=float(len(pipeline_output[pipeline_output["return"]>0]))
m=float(len(pipeline_output[pipeline_output["Trading Signal"]==1]))
a=float(len(pipeline_output[pipeline_output["return"]<0]))
b=float(len(pipeline_output[pipeline_output["Trading Signal"]==0]))
z=float(len(pipeline_output))
print("The percentage of positive returns is:", ((n/z)*100),"%")
print("The percentage of BUY Trading Signal is:", ((m/z)*100),"%")
print("The percentage of negative returns is:", ((a/z)*100),"%")
print("The percentage of SELL Trading Signal is:", ((b/z)*100),"%")
```

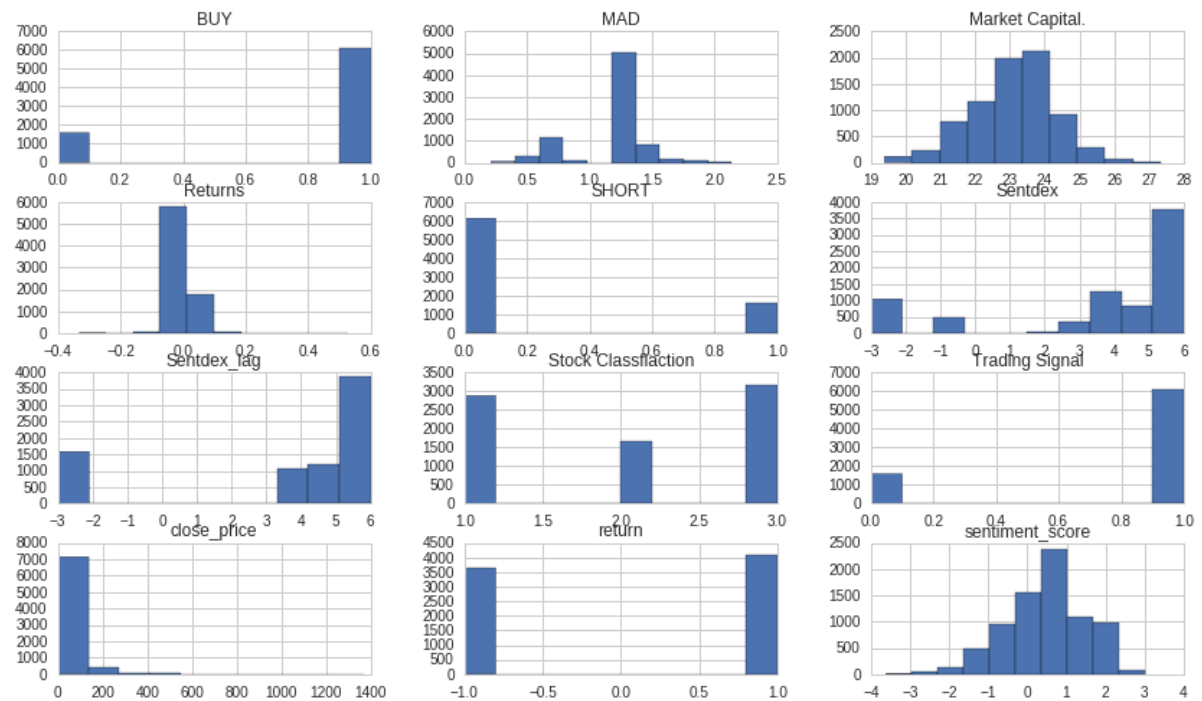
The percentage of positive returns is: 52.7910892371 %
 The percentage of BUY Trading Signal is: 79.406812589 %
 The percentage of negative returns is: 47.2089107629 %
 The percentage of SELL Trading Signal is: 20.593187411 %

```
In [10]: print(pipeline_output['Trading Signal'].value_counts())
print(pipeline_output['return'].value_counts())
```


#Unbalanced Labels and datasets

```
1.0    6131
0.0    1590
Name: Trading Signal, dtype: int64
1.0     4076
-1.0    3645
Name: return, dtype: int64
```

```
In [11]: pipeline_output.hist();
```

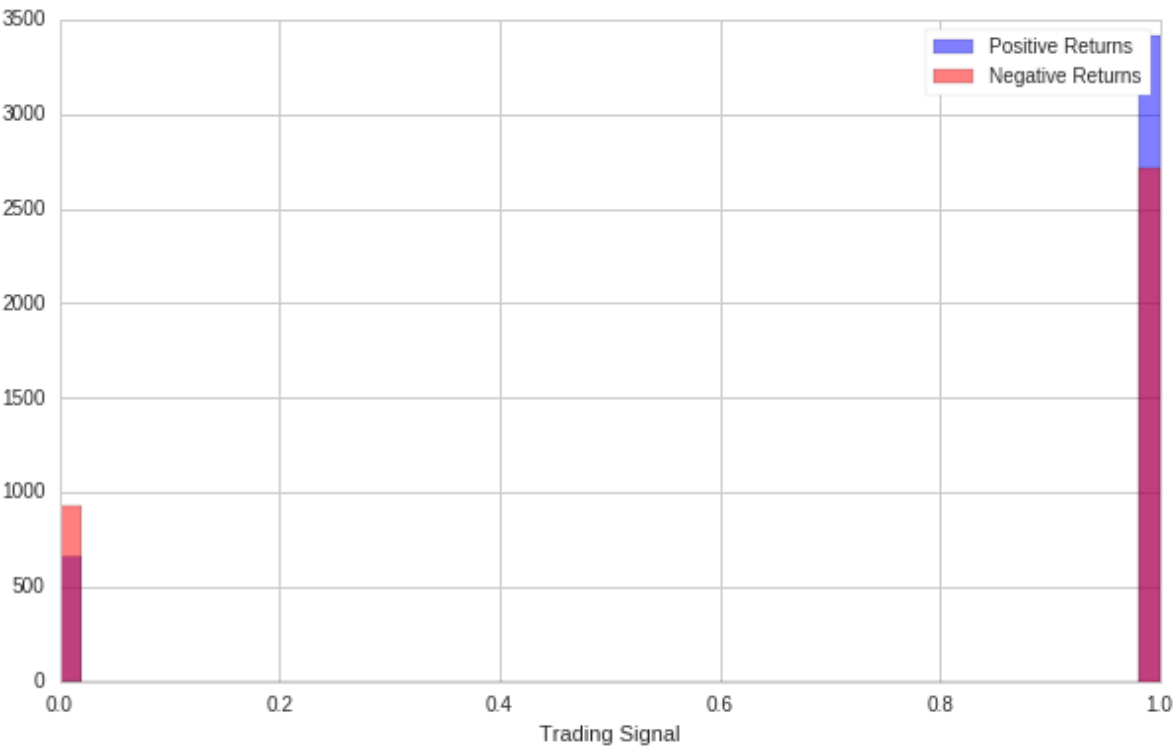



```
In [12]: sns.boxplot(x='Trading Signal',y='Returns',data=pipeline_output,palette='rainbow', width= 0.3);  
# For which I assigned buy (1), what where their returns  
# For which I assigned sell(0), what where their returns
```





```
In [13]: plt.figure(figsize=(10,6))
pipeline_output[pipeline_output['Trading Signal']==1]['Returns'].hist(alpha=0.5,color='blue',
                                                                    bins=50,label='Trading Signal=1'
)
pipeline_output[pipeline_output['Trading Signal']==0]['Returns'].hist(alpha=0.5,color='red',
                                                                    bins=50,label='Trading Signal=0'
)
plt.legend()
plt.xlabel('Returns');

plt.figure(figsize=(10,6))
pipeline_output[pipeline_output['Returns']>0]['Trading Signal'].hist(alpha=0.5,color='blue',
                                                                    bins=50,label='Positive Returns'
)
pipeline_output[pipeline_output['Returns']<0]['Trading Signal'].hist(alpha=0.5,color='red',
                                                                    bins=50,label='Negative Returns'
)
plt.legend()
plt.xlabel('Trading Signal');
```



```
In [14]: plt.figure(figsize=(10,6))
pipeline_output[pipeline_output['Returns']>0]['Stock Classfiaction'].hist(alpha=0.5,color='blue',
                                                                    bins=50,label='Positive return')
pipeline_output[pipeline_output['Returns']<0]['Stock Classfiaction'].hist(alpha=0.5,color='red',
                                                                    bins=50,label='Negative Return')
plt.legend()
plt.xlabel('Stock Classification');



plt.figure(figsize=(10,6))
pipeline_output[pipeline_output['Trading Signal']==1]['Stock Classfiaction'].hist(alpha=0.5,color='blue',
                                                                    bins=50,label='Trading Signal=1')
)
pipeline_output[pipeline_output['Trading Signal']==0]['Stock Classfiaction'].hist(alpha=0.5,color='red',
                                                                    bins=50,label='Trading Signal=0')
)
plt.legend()
plt.xlabel('Stock Classification');
```




```
In [15]: plt.figure(figsize=(10,6))
pipeline_output[pipeline_output['Returns']>0]['Sentdex'].hist(alpha=0.5,color='blue',
                                                             bins=30,label='Positive Returns'
)
pipeline_output[pipeline_output['Returns']<0]['Sentdex'].hist(alpha=0.5,color='red',
                                                             bins=30,label='Negative Returns'
)
plt.legend()
plt.xlabel('Sentdex');

plt.figure(figsize=(10,6))
pipeline_output[pipeline_output['Returns']>0]['Sentdex_lag'].hist(alpha=0.5,color='blue',
                                                             bins=30,label='Positive Returns'
)
pipeline_output[pipeline_output['Returns']<0]['Sentdex_lag'].hist(alpha=0.5,color='red',
                                                             bins=30,label='Negative Returns'
)
plt.legend()
plt.xlabel('Sentdex 7-Day Lag');



plt.figure(figsize=(10,6))
pipeline_output[pipeline_output['Returns']>0]['sentiment_score'].hist(alpha=0.5,color='blue',
                                                             bins=50,label='Positive Returns'
)
pipeline_output[pipeline_output['Returns']<0]['sentiment_score'].hist(alpha=0.5,color='red',
                                                             bins=50,label='Negative Returns'
)
plt.legend()
plt.xlabel('Stocktwits Sentiment');
```



```
In [16]: plt.figure(figsize=(10,6))
pipeline_output[pipeline_output['Trading Signal']==1]['Market Capital.'].hist(
alpha=0.5,color='blue',
bins=30,label='Positive Returns'
)
pipeline_output[pipeline_output['Trading Signal']==0]['Market Capital.'].hist(
alpha=0.5,color='red',
bins=30,label='Negative Returns'
)
plt.legend()
plt.xlabel('Market Capitalization');

plt.figure(figsize=(10,6))
pipeline_output[pipeline_output['Returns']>0]['Market Capital.'].hist(alpha=0.
5,color='blue',
bins=30,label='Positive Returns'
)
pipeline_output[pipeline_output['Returns']<0]['Market Capital.'].hist(alpha=0.
5,color='red',
bins=30,label='Negative Returns'
)
plt.legend()
plt.xlabel('Market Capitalization');
```



```
In [17]: print('Number of securities that passed the filter: %d' % len(pipeline_output.
index.levels[1].unique()))
pipeline_output.columns
```

Number of securities that passed the filter: 365

```
Out[17]: Index([u'BUY', u'MAD', u'Market Capital.', u'Returns', u'SHORT', u'Sentdex',
u'Sentdex_lag', u'Stock Classfiacion', u'close_price', u'return',
u'sentiment_score', u'Trading Signal'],
dtype='object')
```

```
In [18]: pipeline_output.info()
```

```
<class 'pandas.core.frame.DataFrame'>
MultiIndex: 7721 entries, (2013-01-03 00:00:00+00:00, Equity(754 [BBY])) to
(2017-01-03 00:00:00+00:00, Equity(34913 [RF]))
Data columns (total 12 columns):
BUY                7721 non-null bool
MAD                7721 non-null float64
Market Capital.    7721 non-null float64
Returns            7721 non-null float64
SHORT              7721 non-null bool
Sentdex            7721 non-null float64
Sentdex_lag        7721 non-null float64
Stock Classfiacion 7721 non-null int64
close_price        7721 non-null float64
return             7721 non-null float64
sentiment_score    7721 non-null float64
Trading Signal     7721 non-null float64
dtypes: bool(2), float64(9), int64(1)
memory usage: 678.6+ KB
```

****BUILDING THE MODEL PHASE****

```
In [7]: from sklearn.linear_model import LogisticRegression
# Testing how the model's intuition
model = LogisticRegression()
model.fit([[-2,-3],[1,0],[1,1]],["T","F","T"])
# model.coef_, model.intercept_,model.predict([-3,21])
```

```
Out[7]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
intercept_scaling=1, max_iter=100, multi_class='ovr',
penalty='l2', random_state=None, solver='liblinear', tol=0.0001,
verbose=0)
```

```
In [8]: feature_cols = ["Market Capital.", "sentiment_score", "Sentdex_lag", "MAD", "Tra
ding Signal", "Stock Classfiacion", "Sentdex"]
X = pipeline_output[feature_cols]
y = pipeline_output["return"] #target column (classified T or F)
```

```
In [9]: # split X and y into training and testing sets
from sklearn.cross_validation import train_test_split
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.3,random_state=
None)
```

```
In [10]: # instantiate the model (using the default parameters)
logmodel = LogisticRegression()
# fit the model with data
logmodel.fit(X_train,y_train)
```

```
Out[10]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
        intercept_scaling=1, max_iter=100, multi_class='ovr',
        penalty='l2', random_state=None, solver='liblinear', tol=0.0001,
        verbose=0)
```

```
In [11]: #
y_pred=logmodel.predict(X_test)
y_pred
logmodel.coef_, logmodel.intercept_,logmodel.predict(X_test)
```

```
Out[11]: (array([[ 0.05054444,  0.08148023,  0.00740886,  0.18409867,  0.28884548,
        0.00258946, -0.02378607]]),
        array([-1.46557072]),
        array([ 1.,  1., -1., ...,  1.,  1.,  1.]))
```

```
In [12]: # import the metrics class
from sklearn import metrics
cnf_matrix = metrics.confusion_matrix(y_test, y_pred)
cnf_matrix
```

```
Out[12]: array([[ 416, 1503],
        [ 297, 1885]])
```

```
In [13]: print"Accuracy:",metrics.accuracy_score(y_test, y_pred)
print"Precision:",metrics.precision_score(y_test, y_pred)
print"Recall:",metrics.recall_score(y_test, y_pred)
```

```
Accuracy: 0.561082662765
Precision: 0.556375442739
Recall: 0.863886342805
```

```
In [14]: from sklearn.metrics import classification_report
print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
-1.0	0.58	0.22	0.32	1919
1.0	0.56	0.86	0.68	2182
avg / total	0.57	0.56	0.51	4101

Determine Threshold

```
In [23]: from sklearn.metrics import accuracy_score, confusion_matrix, recall_score, roc_auc_score, precision_score


THRESHOLD = 0.50
preds = np.where(logmodel.predict_proba(X_test)[:,-1] > THRESHOLD, 1, -1)

pd.DataFrame(data=[accuracy_score(y_test, preds), recall_score(y_test, preds),
                      precision_score(y_test, preds), roc_auc_score(y_test, preds)],
              index=["accuracy", "recall", "precision", "roc_auc_score"], columns = ["Scores"])
# 0.5 remained the best threshold
```

Out[23]:


	Scores
accuracy	0.564091
recall	0.851485
precision	0.554243
roc_auc_score	0.550177

```
In [24]: class_names=[-1,1] # name of classes
fig, ax = plt.subplots()
tick_marks = np.arange(len(class_names))
plt.xticks(tick_marks, class_names)
plt.yticks(tick_marks, class_names)
# create heatmap
sns.heatmap(pd.DataFrame(cnf_matrix), annot=True, cmap="YlGnBu", fmt='g')
ax.xaxis.set_label_position("top")
plt.tight_layout()
plt.title('Confusion matrix', y=1.1)
plt.ylabel('Actual label')
plt.xlabel('Predicted label');
```



AUROC represents the likelihood of the model distinguishing observations from two classes. In other words, if a random selection of an observation from each class is made, what's the probability that your model will be able to "rank" them correctly?


```
In [25]: y_pred_proba = logmodel.predict_proba(X_test)[::,1]
fpr, tpr, threshold = metrics.roc_curve(y_test, y_pred_proba)
auc = metrics.roc_auc_score(y_test, y_pred_proba)
plt.plot(fpr,tpr,label="data 1, auc="+str(auc))
x = np.linspace(0, 1, 100000)
plt.plot(x, x + 0, linestyle='solid',c = 'k')
plt.legend(loc=4)
plt.xlabel("FPR")
plt.ylabel("TPR")
plt.show()
```



Random Forest Classifier

```
In [15]: from sklearn.ensemble import RandomForestClassifier

rfc = RandomForestClassifier(n_estimators=800, max_leaf_nodes=2,max_features=7
, min_samples_split=2)
rfc.fit(X_train, y_train)
```

```
Out[15]: RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
max_depth=None, max_features=7, max_leaf_nodes=2,
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, n_estimators=800, n_jobs=1,
oob_score=False, random_state=None, verbose=0,
warm_start=False)
```

```
In [16]: from sklearn import metrics
rfc_pred = rfc.predict(X_test)
rfc_pred
cnf_matrix = metrics.confusion_matrix(y_test, rfc_pred)
cnf_matrix
```

```
Out[16]: array([[ 423, 1496],
               [ 314, 1868]])
```


```
In [17]: print("Accuracy:",metrics.accuracy_score(y_test, rfc_pred))
print("Precision:",metrics.precision_score(y_test, rfc_pred))
print("Recall:",metrics.recall_score(y_test, rfc_pred))
```

```
('Accuracy:', 0.57449402584735432)
('Precision:', 0.57298731937481573)
('Recall:', 0.86741071428571426)
```

```
In [18]: print(classification_report(y_test,rfc_pred))
```

	precision	recall	f1-score	support
-1.0	0.58	0.22	0.32	1861
1.0	0.57	0.87	0.69	2240
avg / total	0.58	0.57	0.52	4101

```
In [41]: y_pred_proba = rfc.predict_proba(X_test)[::,1]
fpr, tpr, _ = metrics.roc_curve(y_test, y_pred_proba)
auc = metrics.roc_auc_score(y_test, y_pred_proba)
plt.plot(fpr,tpr,label="data 1, auc="+str(auc))
x = np.linspace(0, 1, 100000)
plt.plot(x, x + 0, linestyle='solid',c = 'k')
plt.legend(loc=4)
plt.show()
```



Support Vector Machine (SVMs)

```
In [17]: from sklearn.svm import SVC
sv_model = SVC(C=1,gamma=0.1)
sv_model.fit(X_train,y_train)
```

```
Out[17]: SVC(C=1, cache_size=200, class_weight=None, coef0=0.0, degree=3, gamma=0.1,
kernel='rbf', max_iter=-1, probability=False, random_state=None,
shrinking=True, tol=0.001, verbose=False)
```

```
In [22]: sv_predictions = sv_model.predict(X_test)
print(metrics.confusion_matrix(y_test,sv_predictions))
print("\n")
print(classification_report(y_test,sv_predictions))
```

```
[[ 411 1539]
 [ 306 1845]]
```

	precision	recall	f1-score	support
-1.0	0.57	0.21	0.31	1950
1.0	0.55	0.86	0.67	2151
avg / total	0.56	0.55	0.50	4101

```
In [22]: ##### GRIDSEARCH #####
#Hypertuning parameters
param_grid = {'C': [0.1,1, 10, 100, 1000], 'gamma': [1,0.1,0.01,0.001,0.0001]}
from sklearn.grid_search import GridSearchCV
```

```
In [23]: grid = GridSearchCV(SVC(),param_grid,refit=True,verbose=0)
```

```
In [24]: grid.fit(X_train,y_train)
```

```
Out[24]: GridSearchCV(cv=None, error_score='raise',
      estimator=SVC(C=1.0, cache_size=200, class_weight=None, coef0=0.0, degree=3, gamma=0.0,
      kernel='rbf', max_iter=-1, probability=False, random_state=None,
      shrinking=True, tol=0.001, verbose=False),
      fit_params={}, iid=True, loss_func=None, n_jobs=1,
      param_grid={'C': [0.1, 1, 10, 100, 1000], 'gamma': [1, 0.1, 0.01, 0.001, 0.0001]},
      pre_dispatch='2*n_jobs', refit=True, score_func=None, scoring=None,
      verbose=0)
```

```
In [25]: print(grid.best_params_)
grid.best_estimator_

{'C': 1, 'gamma': 0.1}
```

```
Out[25]: SVC(C=1, cache_size=200, class_weight=None, coef0=0.0, degree=3, gamma=0.1,
      kernel='rbf', max_iter=-1, probability=False, random_state=None,
      shrinking=True, tol=0.001, verbose=False)
```

```
In [26]: grid_predictions = grid.predict(X_test)
print(metrics.confusion_matrix(y_test,grid_predictions))
print("\n")
print(classification_report(y_test,grid_predictions))
```

```
[[ 298  769]
 [ 229 1021]]
```

	precision	recall	f1-score	support
-1.0	0.57	0.28	0.37	1067
1.0	0.57	0.82	0.67	1250
avg / total	0.57	0.57	0.53	2317

K-Nearest Neighbor

Run it after Pipeliene directly

```
In [51]: from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
```

```
In [52]: scaler.fit(pipeline_output.drop(['close_price', 'BUY', 'return', 'Returns', 'S
HORT'],axis=1))
```

```
Out[52]: StandardScaler(copy=True, with_mean=True, with_std=True)
```

```
In [53]: scaled_features = scaler.transform(pipeline_output.drop(['close_price', 'BUY',
'return', 'Returns', 'SHORT'],axis=1))
```

```
In [54]: scaled_features
```

```
Out[54]: array([[ -1.92110916, -0.8276163 , -2.12216652, ..., -1.17211062,
-1.06454103, -1.96366363],
[ -1.92306567, -0.8276163 , -1.49599865, ..., -1.17211062,
-1.06001677, -1.96366363],
[  2.57138894, -2.04815612, -0.24366291, ..., -1.17211062,
-0.53195792,  0.50925219],
...,
[  0.3088213 ,  1.44250233,  0.69558889, ..., -1.17211062,
 0.76959902,  0.50925219],
[  0.42621986,  1.25449715,  0.06942102, ...,  1.09063874,
 0.22452827,  0.50925219],
[  1.15531608,  0.42339293,  0.38250495, ..., -1.17211062,
 0.50079701,  0.50925219]])
```

```
In [56]: pipeline_scaled = pd.DataFrame(scaled_features, columns=pipeline_output.column
s.drop(['close_price', 'BUY', 'return', 'Returns', 'SHORT']))
```

In [57]: `pipeline_scaled.head()`

Out[57]:

	MAD	Market Capital.	Sentdex	Sentdex_lag	Stock Classfiation	sentiment_score	Trading Signal
0	-1.921109	-0.827616	-2.122167	-1.925463	-1.172111	-1.064541	-1.963664
1	-1.923066	-0.827616	-1.495999	-1.925463	-1.172111	-1.060017	-1.963664
2	2.571389	-2.048156	-0.243663	0.654283	-1.172111	-0.531958	0.509252
3	-1.917113	-0.827616	-1.495999	-1.925463	-1.172111	-1.825834	-1.963664
4	0.494555	1.392060	0.695589	0.654283	-0.040736	0.024551	0.509252

In [49]: `from sklearn.cross_validation import train_test_split`
`X_train, X_test, y_train, y_test = train_test_split(scaled_features, pipeline_o`
`utput['return'],`
`test_size=0.30)`
Get dummies for returns first before running

In [70]: `from sklearn.neighbors import KNeighborsClassifier`
`knn = KNeighborsClassifier(n_neighbors=36)`

In [71]: `knn.fit(X_train, y_train)`

Out[71]: `KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',`
`metric_params=None, n_neighbors=36, p=2, weights='uniform')`

In [72]: `pred = knn.predict(X_test)`
`from sklearn.metrics import classification_report, confusion_matrix`
`print(confusion_matrix(y_test, pred))`
`print("\n")`
`print(classification_report(y_test, pred))`

```
[[503 602]
 [455 757]]
```

	precision	recall	f1-score	support
-1.0	0.53	0.46	0.49	1105
1.0	0.56	0.62	0.59	1212
avg / total	0.54	0.54	0.54	2317

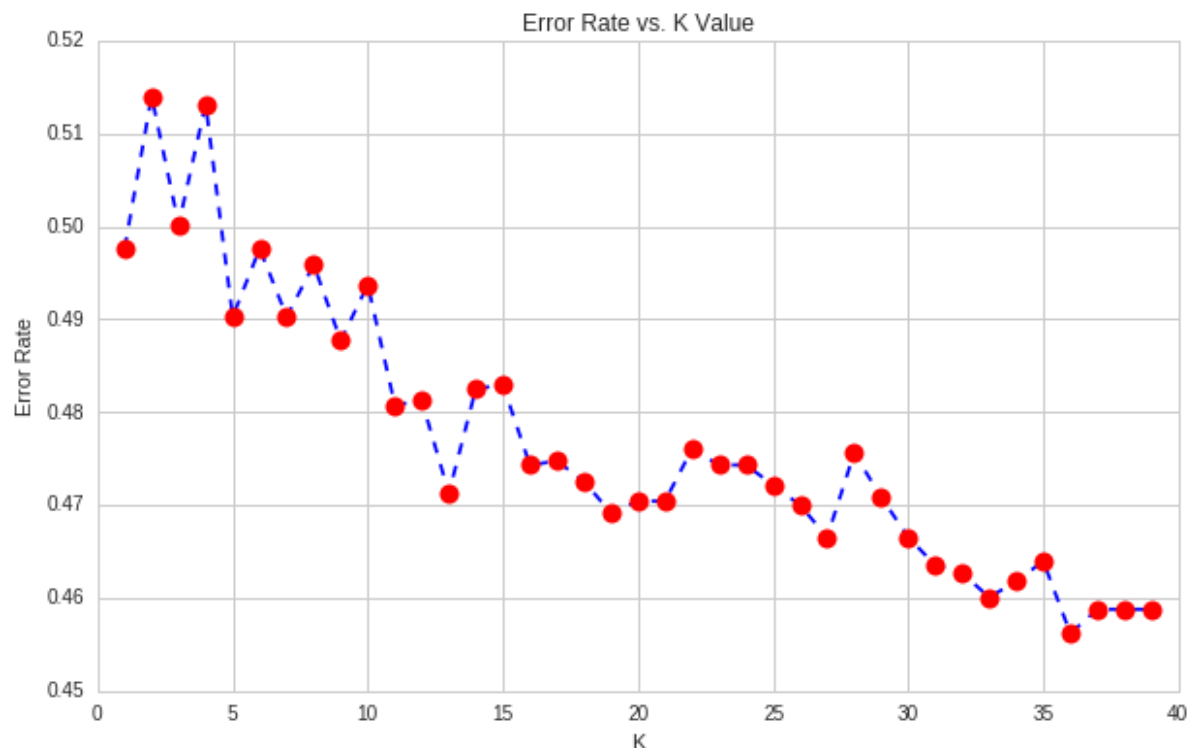
```
In [62]: error_rate = []

# Will take some time
for i in range(1,40):

    knn = KNeighborsClassifier(n_neighbors=i)
    knn.fit(X_train,y_train)
    pred_i = knn.predict(X_test)
    error_rate.append(np.mean(pred_i != y_test))
```

```
In [63]: plt.figure(figsize=(10,6))
plt.plot(range(1,40),error_rate,color='blue', linestyle='dashed', marker='o',
         markerfacecolor='red', markersize=10)
plt.title('Error Rate vs. K Value')
plt.xlabel('K')
plt.ylabel('Error Rate')
```

Out[63]: <matplotlib.text.Text at 0x7f1754058790>



```
In [ ]: from sklearn.linear_model import LinearRegression
lm = LinearRegression()
lm.fit(X_train,y_train)
coeff_df = pd.DataFrame(lm.coef_,X.columns,columns=['Coefficient'])
coeff_df
predictions = lm.predict(X_test)
plt.scatter(y_test,predictions)
plt.xlabel('Y Test')
plt.ylabel('Predicted Y')
from sklearn import metrics
print('MAE:', metrics.mean_absolute_error(y_test, predictions))
print('MSE:', metrics.mean_squared_error(y_test, predictions))
print('RMSE:', np.sqrt(metrics.mean_squared_error(y_test, predictions)))
```

Out of Sample Predictions:


```
In [74]: pipeline_outsample = run_pipeline(  
        make_pipeline(),  
        start_date="2019-01-01",  
        end_date="2019-06-21"  
    )  
    pipeline_outsample.tail(20)
```

Pipeline Execution Time: 2 Minutes, 18.30 Seconds

Out[74]:

		BUY	MAD	Market Capital.	Returns	SHORT	Sentdex	Sentdex
2019-01-02 00:00:00+00:00	Equity(3149 [GE])	False	0.571820	6.584473e+10	0.005984	True	-3.0	
	Equity(4705 [MKC])	True	1.234640	1.832881e+10	0.002015	False	5.0	
	Equity(5029 [MRK])	True	1.202823	1.986948e+11	0.014727	False	6.0	
	Equity(13635 [DO])	False	0.634035	1.297384e+09	-0.031828	True	-3.0	
	Equity(337 [AMAT])	False	0.723636	3.209413e+10	0.022283	True	2.0	
2019-01-03 00:00:00+00:00	Equity(7447 [TIF])	False	0.742030	9.974845e+09	0.016522	True	-1.0	
	Equity(13635 [DO])	False	0.626310	1.338614e+09	0.032874	True	-3.0	
	Equity(26143 [NRG])	True	1.206547	1.104923e+10	-0.037626	False	6.0	
	Equity(51649 [ADT])	False	0.783810	4.915695e+09	0.064007	True	-3.0	
	Equity(337 [AMAT])	False	0.718949	3.023444e+10	-0.058525	True	6.0	
2019-01-04 00:00:00+00:00	Equity(42950 [FB])	False	0.799401	3.785928e+11	-0.026828	True	-3.0	
	Equity(4705 [MKC])	True	1.213979	1.798623e+10	0.002936	False	5.0	
	Equity(9883 [ATVI])	False	0.681487	3.599310e+10	0.040132	True	-1.0	
	Equity(12909 [LH])	False	0.777963	1.291116e+10	0.034196	True	4.0	
	Equity(23269 [WW])	False	0.631015	2.494360e+09	0.021067	True	-3.0	
2019-01-07 00:00:00+00:00	Equity(32660 [SFLY])	False	0.573869	1.460114e+09	0.054598	True	-1.0	
	Equity(42950 [FB])	False	0.799374	3.964390e+11	0.044532	True	-3.0	
	Equity(4705 [MKC])	True	1.207199	1.815031e+10	0.008708	False	5.0	
	Equity(12909 [LH])	False	0.774010	1.281228e+10	-0.007504	True	2.0	
	Equity(23269 [WW])	False	0.621930	2.402793e+09	-0.037245	True	-3.0	

In [75]: `pipeline_outsample['return'] = pipeline_outsample["return"].apply(np.sign)`

```
In [76]: pipeline_outsample["Trading Signal"] = pd.get_dummies(pipeline_outsample['BUY'],drop_first=True)
```

```
In [77]: pipeline_outsample['Market Capital.']= pipeline_outsample["Market Capital."].
         apply(np.log)

pipeline_outsample.head(20)
```

Out[77]:

		BUY	MAD	Market Capital.	Returns	SHORT	Sentdex	Sentdex_k
2019-01-02 00:00:00+00:00	Equity(3149 [GE])	False	0.571820	24.910565	0.005984	True	-3.0	-3
	Equity(4705 [MKC])	True	1.234640	23.631740	0.002015	False	5.0	5
	Equity(5029 [MRK])	True	1.202823	26.015036	0.014727	False	6.0	6
	Equity(13635 [DO])	False	0.634035	20.983615	-0.031828	True	-3.0	-3
	Equity(337 [AMAT])	False	0.723636	24.191939	0.022283	True	2.0	-3
2019-01-03 00:00:00+00:00	Equity(7447 [TIF])	False	0.742030	23.023332	0.016522	True	-1.0	-3
	Equity(13635 [DO])	False	0.626310	21.014901	0.032874	True	-3.0	-3
	Equity(26143 [NRG])	True	1.206547	23.125627	-0.037626	False	6.0	6
	Equity(51649 [ADT])	False	0.783810	22.315699	0.064007	True	-3.0	-3
	Equity(337 [AMAT])	False	0.718949	24.132247	-0.058525	True	6.0	-3
2019-01-04 00:00:00+00:00	Equity(42950 [FB])	False	0.799401	26.659727	-0.026828	True	-3.0	-3
	Equity(4705 [MKC])	True	1.213979	23.612872	0.002936	False	5.0	5
	Equity(9883 [ATVI])	False	0.681487	24.306593	0.040132	True	-1.0	-3
	Equity(12909 [LH])	False	0.777963	23.281358	0.034196	True	4.0	-3
	Equity(23269 [WW])	False	0.631015	21.637298	0.021067	True	-3.0	-3
2019-01-07 00:00:00+00:00	Equity(32660 [SFLY])	False	0.573869	21.101781	0.054598	True	-1.0	-3
	Equity(42950 [FB])	False	0.799374	26.705788	0.044532	True	-3.0	-3
	Equity(4705 [MKC])	True	1.207199	23.621954	0.008708	False	5.0	5
	Equity(12909 [LH])	False	0.774010	23.273670	-0.007504	True	2.0	-3
	Equity(23269 [WW])	False	0.621930	21.599898	-0.037245	True	-3.0	-3

```
In [78]: X_ofs17 = pipeline_outsample[feature_cols]
y_ofs17 = pipeline_outsample["return"]
len(X_ofs17)
```

Out[78]: 525

Fitting Logistic Model to out of sample data:

```
In [79]: y_pred_ofs17 = logmodel.predict(X_ofs17)
y_pred_ofs17
```

Out[79]: array([-1., 1., 1., -1., -1., -1., -1., 1., -1., -1., -1., 1., -1.,
-1., -1., -1., -1., 1., -1., -1., -1., -1., -1., -1., -1., -1.,
-1., -1., -1., -1., -1., -1., -1., -1., -1., -1., -1., -1., -1.,
-1., -1., 1., -1., 1., -1., -1., -1., -1., -1., -1., -1., -1.,
-1., -1., -1., -1., -1., -1., -1., -1., -1., -1., -1., -1., 1.,
-1., -1., -1., -1., -1., -1., -1., -1., -1., 1., -1., -1., -1.,
-1., -1., -1., -1., -1., -1., -1., -1., -1., -1., -1., -1., -1.,
-1., -1., -1., -1., -1., -1., -1., -1., -1., 1., 1., -1., -1.,
1., -1., -1., -1., -1., -1., -1., 1., -1., 1., -1., -1., -1.,
-1., -1., -1., -1., -1., -1., -1., -1., -1., 1., -1., -1., -1.,
1., 1., 1., -1., 1., 1., 1., -1., 1., -1., 1., 1.,
1., -1., 1., 1., 1., 1., 1., 1., -1., 1., -1., 1., 1.,
-1., -1., 1., -1., 1., -1., 1., 1., -1., 1., -1., 1., -1.,
1., -1., 1., 1., 1., 1., 1., 1., -1., 1., -1., 1., 1.,
1., 1., 1., -1., 1., 1., 1., 1., 1., 1., -1., 1., 1.,
1., -1., 1., 1., -1., 1., 1., 1., -1., 1., 1., -1., 1.,
-1., 1., 1., 1., 1., -1., 1., 1., -1., 1., 1., 1., 1.,
1., -1., 1., 1., -1., 1., 1., -1., 1., 1., 1., -1., 1.,
1., -1., 1., 1., -1., 1., -1., -1., 1., 1., -1., 1., -1.,
1., 1., 1., -1., -1., -1., 1., 1., -1., -1., 1., 1.,
-1., -1., 1., 1., -1., -1., -1., 1., 1., 1., -1., -1., -1.,
1., -1., 1., -1., -1., 1., 1., 1., 1., -1., 1., 1., 1.,
1., 1., 1., 1., 1., 1., 1., -1., 1., 1., 1., 1., 1.,
1., -1., 1., -1., 1., -1., -1., 1., 1., 1., -1., 1., -1.,
-1., 1., 1., -1., 1., -1., 1., 1., -1., 1., 1., 1., -1.,
1., 1., 1., -1., 1., 1., 1., 1., 1., 1., -1., 1., 1.,
-1., -1., 1., -1., -1., -1., 1., -1., -1., 1., -1., -1.,
-1., 1., 1., -1., -1., -1., -1., 1., 1., -1., 1., -1.,
1., 1., 1., -1., -1., -1., 1., 1., -1., 1., 1., 1., -1.,
1., -1., -1., 1., -1., -1., -1., 1., 1., 1., -1., -1.,
1., 1., 1., -1., 1., -1., -1., 1., -1., -1., 1., 1., 1.,
-1., 1., 1., -1., 1., 1., 1., 1., -1., -1., 1., 1., 1.,
-1., 1., 1., -1., 1., 1., -1., 1., 1., -1., 1., 1., 1.,
1., -1., 1., -1., 1.,

```
In [80]: cnf_matrix = metrics.confusion_matrix(y_ofs17, y_pred_ofs17)
print cnf_matrix
print "Accuracy:", metrics.accuracy_score(y_ofs17, y_pred_ofs17)
print "Precision:", metrics.precision_score(y_ofs17, y_pred_ofs17)
print "Recall:", metrics.recall_score(y_ofs17, y_pred_ofs17)
print classification_report(y_ofs17, y_pred_ofs17)
```

```
[[132  93]
 [143 157]]
Accuracy: 0.550476190476
Precision: 0.628
Recall: 0.523333333333
```

	precision	recall	f1-score	support
-1.0	0.48	0.59	0.53	225
1.0	0.63	0.52	0.57	300
avg / total	0.56	0.55	0.55	525

In []:

Fitting Random Forest Classifier to out of sample data:

```
In [81]: y_rfc_pred_ofs_17 = rfc.predict(X_ofs17)
         y_rfc_pred_ofs_17
```

```
Out[81]: array([-1.,  1.,  1., -1., -1., -1., -1.,  1., -1., -1., -1.,  1., -1.,
        -1.,  1., -1.,  1.,  1.,  1., -1., -1., -1., -1., -1.,  1.,
        -1., -1.,  1., -1., -1., -1.,  1., -1., -1., -1.,  1., -1.,  1.,
        -1.,  1.,  1.,  1.,  1., -1.,  1., -1.,  1., -1., -1., -1., -1.,
        -1., -1., -1., -1., -1.,  1., -1., -1., -1., -1.,  1., -1.,  1.,
        -1., -1.,  1., -1., -1., -1., -1., -1., -1., -1.,  1., -1., -1.,
        -1., -1., -1., -1., -1., -1., -1., -1., -1.,  1.,  1., -1., -1.,
         1.,  1., -1., -1., -1.,  1., -1.,  1., -1.,  1., -1., -1., -1.,
        -1., -1., -1., -1., -1., -1., -1., -1., -1., -1.,  1., -1., -1.,
         1.,  1.,  1., -1.,  1.,  1.,  1., -1.,  1., -1., -1.,  1., -1.,
         1., -1.,  1.,  1.,  1.,  1.,  1.,  1.,  1.,  1., -1.,  1.,  1.,
        -1.,  1.,  1., -1.,  1.,  1.,  1.,  1.,  1.,  1., -1.,  1.,  1.,
         1., -1.,  1.,  1.,  1.,  1.,  1.,  1.,  1.,  1., -1.,  1.,  1.,
        -1.,  1.,  1.,  1.,  1., -1.,  1.,  1., -1.,  1.,  1.,  1.,  1.,
         1., -1.,  1.,  1., -1., -1., -1., -1., -1.,  1., -1.,  1., -1.,
         1.,  1.,  1., -1., -1., -1.,  1.,  1.,  1., -1., -1., -1., -1.,
        -1., -1.,  1.,  1., -1., -1., -1.,  1.,  1.,  1., -1., -1., -1.,
         1., -1.,  1., -1., -1.,  1.,  1.,  1.,  1., -1.,  1.,  1.,  1.,
         1.,  1.,  1., -1.,  1.,  1.,  1., -1.,  1.,  1.,  1., -1.,  1.,
         1., -1.,  1., -1.,  1., -1., -1.,  1.,  1.,  1., -1.,  1.,  1.,
        -1.,  1.,  1., -1.,  1., -1.,  1.,  1., -1.,  1.,  1.,  1., -1.,
         1.,  1.,  1., -1.,  1.,  1.,  1.,  1.,  1., -1., -1.,  1.,  1.,
        -1.,  1., -1.,  1., -1., -1., -1.,  1.,  1.,  1., -1., -1., -1.,
         1., -1., -1.,  1., -1., -1., -1.,  1.,  1.,  1., -1., -1., -1.,
         1.,  1.,  1., -1.,  1., -1., -1.,  1., -1., -1.,  1.,  1.,  1.,
        -1.,  1.,  1., -1.,  1.,  1.,  1.,  1., -1., -1.,  1.,  1.,  1.,
        -1.,  1.,  1., -1.,  1.,  1., -1.,  1.,  1., -1.,  1.,  1.,  1.,
         1., -1.,  1., -1.,  1.]
```

```
In [82]: cnf_matrix_rfc = metrics.confusion_matrix(y_ofs17, y_rfc_pred_ofs_17)
print(cnf_matrix_rfc)
print("Accuracy:", metrics.accuracy_score(y_ofs17, y_rfc_pred_ofs_17))
print("Precision:", metrics.precision_score(y_ofs17, y_rfc_pred_ofs_17))
print("Recall:", metrics.recall_score(y_ofs17, y_rfc_pred_ofs_17))
print(classification_report(y_ofs17, y_rfc_pred_ofs_17))
```

```
[[121 104]
 [129 171]]
('Accuracy:', 0.55619047619047624)
('Precision:', 0.62181818181818183)
('Recall:', 0.56999999999999995)
      precision    recall  f1-score   support

   -1.0         0.48         0.54         0.51         225
    1.0         0.62         0.57         0.59         300

avg / total         0.56         0.56         0.56         525
```

```
In [ ]: ### Fitting Support Vector Machines
```



```
In [84]: y_svm_pred_ofs_17 = sv_model.predict(X_ofs17)
y_svm_pred_ofs_17
```

```
Out[84]: array([-1.,  1.,  1., -1.,  1., -1., -1.,  1., -1., -1., -1.,  1., -1.,
-1., -1., -1., -1.,  1.,  1., -1., -1., -1., -1., -1., -1.,
-1., -1., -1., -1.,  1., -1., -1., -1., -1., -1., -1., -1., -1.,
-1., -1.,  1., -1.,  1., -1., -1., -1., -1., -1.,  1., -1., -1.,
-1., -1., -1., -1., -1., -1., -1., -1., -1., -1., -1.,  1.,  1.,
-1., -1., -1., -1., -1., -1., -1., -1., -1.,  1., -1., -1., -1.,
-1., -1., -1., -1., -1., -1., -1., -1., -1.,  1.,  1., -1., -1.,
  1., -1., -1., -1., -1., -1., -1.,  1., -1.,  1., -1., -1., -1.,
-1., -1., -1., -1., -1., -1., -1., -1., -1.,  1., -1., -1., -1.,
  1.,  1.,  1., -1.,  1.,  1.,  1., -1.,  1., -1., -1.,  1., -1.,
  1., -1.,  1.,  1.,  1.,  1.,  1.,  1., -1.,  1., -1.,  1.,  1.,
-1., -1.,  1., -1.,  1., -1.,  1.,  1., -1.,  1., -1.,  1., -1.,
  1., -1.,  1., -1.,  1., -1.,  1.,  1.,  1., -1.,  1.,  1., -1.,
  1.,  1.,  1.,  1.,  1., -1.,  1.,  1.,  1., -1.,  1.,  1., -1.,
  1., -1.,  1., -1.,  1., -1., -1.,  1.,  1.,  1., -1.,  1., -1.,
-1.,  1.,  1.,  1.,  1., -1.,  1.,  1., -1.,  1.,  1.,  1.,  1.,
  1., -1.,  1., -1.,  1.,  1.,  1.,  1.,  1., -1.,  1., -1.,  1.,
-1., -1.,  1., -1., -1., -1.,  1., -1., -1.,  1., -1., -1., -1.,
-1.,  1.,  1., -1., -1., -1., -1.,  1.,  1., -1.,  1., -1., -1.,
  1.,  1., -1.,  1., -1., -1., -1.,  1.,  1.,  1., -1.,  1., -1.,
  1.,  1.,  1., -1., -1., -1.,  1.,  1., -1.,  1.,  1.,  1.,
-1.,  1.,  1., -1.,  1.,  1.,  1.,  1., -1., -1.,  1.,  1.,  1.,
-1.,  1.,  1., -1.,  1.,  1., -1.,  1.,  1., -1., -1.,  1.,  1.,
  1., -1.,  1., -1.,  1.]
```

```
In [86]: cnf_matrix_svm = metrics.confusion_matrix(y_ofs17, y_svm_pred_ofs_17)
print(cnf_matrix_svm)
print("Accuracy:", metrics.accuracy_score(y_ofs17, y_svm_pred_ofs_17))
print("Precision:", metrics.precision_score(y_ofs17, y_svm_pred_ofs_17))
print("Recall:", metrics.recall_score(y_ofs17, y_svm_pred_ofs_17))
print(classification_report(y_ofs17, y_svm_pred_ofs_17))
```

```
[[129  96]
 [141 159]]
('Accuracy:', 0.5485714285714286)
('Precision:', 0.62352941176470589)
('Recall:', 0.53000000000000003)
      precision    recall  f1-score   support

    -1.0         0.48         0.57         0.52         225
     1.0         0.62         0.53         0.57         300

avg / total         0.56         0.55         0.55         525
```

In [87]: `pipeline_outsample.tail(20)`

Out[87]:

		BUY	MAD	Market Capital.	Returns	SHORT	Sentdex	Sentdex_k
2019-06-19 00:00:00+00:00	Equity(460 [APD])	True	1.240573	24.598447	0.009592	False	6.0	6
	Equity(4974 [MSI])	True	1.204562	24.026651	0.007507	False	6.0	6
	Equity(5773 [PBI])	False	0.629607	20.428344	-0.021378	True	-3.0	-3
	Equity(8352 [XRAY])	True	1.272900	23.310976	0.003354	False	6.0	4
	Equity(18221 [VRSN])	True	1.207384	23.923710	0.012737	False	5.0	5
2019-06-20 00:00:00+00:00	Equity(42270 [CPRI])	False	0.679690	22.378551	0.023311	True	-3.0	-3
	Equity(460 [APD])	True	1.242997	24.605139	0.006578	False	6.0	6
	Equity(4974 [MSI])	True	1.210011	24.035395	0.008541	False	6.0	6
	Equity(5773 [PBI])	False	0.627407	20.442802	0.014563	True	-3.0	-3
	Equity(8352 [XRAY])	True	1.275001	23.326859	0.016010	False	6.0	4
2019-06-21 00:00:00+00:00	Equity(18221 [VRSN])	True	1.209886	23.940468	0.016413	False	5.0	5
	Equity(42270 [CPRI])	False	0.676282	22.365783	-0.012399	True	-3.0	-3
	Equity(460 [APD])	True	1.245247	24.612643	0.007624	False	6.0	6
	Equity(939 [BLL])	True	1.290730	23.837294	0.009002	False	6.0	4
	Equity(3676 [HRS])	True	1.209753	23.882572	0.005192	False	4.0	6
2019-06-21 00:00:00+00:00	Equity(4974 [MSI])	True	1.214935	24.037315	0.002823	False	6.0	6
	Equity(5773 [PBI])	False	0.623981	20.413674	-0.028708	True	-3.0	-3
	Equity(8352 [XRAY])	True	1.276252	23.321302	-0.005714	False	6.0	4
	Equity(42270 [CPRI])	False	0.672756	22.386306	0.020146	True	-3.0	-3
	Equity(42277 [ZNGA])	True	1.357073	22.459718	0.002473	False	4.0	4

```
In [124]: pipeline_outsample.iloc[pipeline_outsample.index.levels[0] == '2019-06-21 00:00:00+00:00']
```

```
Out[124]:
```

		BUY	MAD	Market Capital.	Returns	SHORT	Sentdex	Sentdex_lag
2019-03-04 00:00:00+00:00	Equity(939 [BLL])	True	1.253408	23.638807	0.008216	False	6.0	4.0

Stepping Up a year and refitting both models:

```
In [44]: pipeline_output_2 = run_pipeline(
    make_pipeline(),
    start_date="2014-01-01",
    end_date="2018-01-01"
)
```

```
In [45]: pipeline_output_2['return'] = pipeline_output_2["return"].apply(np.sign)
pipeline_output_2['Market Capital.'] = pipeline_output_2["Market Capital."].ap
ply(np.log)
pipeline_output_2["Trading Signal"] = pd.get_dummies(pipeline_output_2['BUY'],
drop_first=True)
```

```
In [46]: n=float(len(pipeline_output_2[pipeline_output_2["return"]>0]))
m=float(len(pipeline_output_2[pipeline_output_2["Trading Signal"]==1]))
a=float(len(pipeline_output_2[pipeline_output_2["return"]<0]))
b=float(len(pipeline_output_2[pipeline_output_2["Trading Signal"]==0]))
x=float(len(pipeline_output_2[pipeline_output_2["Returns"]>0]))
y=float(len(pipeline_output_2[pipeline_output_2["Returns"]<0]))
z=float(len(pipeline_output_2))
print("The percentage of positive returns is:", ((n/z)*100), "%")
print("The percentage of BUY Trading Signal is:", ((m/z)*100), "%")
print("The percentage of negative returns is:", ((a/z)*100), "%")
print("The percentage of SELL Trading Signal is:", ((b/z)*100), "%")
```

The percentage of positive returns is: 46.2763890688 %
The percentage of BUY Trading Signal is: 63.644605621 %
The percentage of negative returns is: 43.970988214 %
The percentage of SELL Trading Signal is: 26.6027716617 %

In [47]: pipeline_output_2.tail(20)

Out[47]:

		BUY	MAD	Market Capital.	Returns	SHORT	Sentdex	Sentdex_la
2017-12-29 00:00:00+00:00	Equity(24124 [WYNN])	True	1.279001	23.585159	-0.010077	False	4.0	6.
	Equity(24811 [GES])	True	1.273858	21.078006	-0.011541	False	6.0	6.
	Equity(24832 [RL])	True	1.201872	22.851096	-0.002140	False	4.0	4.
	Equity(25920 [MAR])	True	1.273396	24.628752	0.002422	False	6.0	5.
	Equity(27676 [AMP])	True	1.234831	23.953562	0.004271	False	4.0	4.
	Equity(32902 [FSLR])	True	1.555720	22.695841	-0.006538	False	6.0	4.
	Equity(41636 [MPC])	True	1.214153	24.202668	-0.000301	False	6.0	6.
	Equity(114 [ADBE])	True	1.215385	25.178645	-0.001424	False	6.0	6.
	Equity(1267 [CAT])	True	1.331272	25.263894	-0.005301	False	6.0	6.
	Equity(1539 [CI])	True	1.202831	24.636723	-0.010037	False	5.0	5.
	Equity(1941 [CTAS])	True	1.204947	23.532140	-0.004980	False	4.0	4.
	Equity(3321 [GPS])	True	1.343680	23.306846	-0.010456	False	5.0	5.
	Equity(6546 [ROST])	True	1.221500	24.146460	-0.004094	False	6.0	6.
	Equity(7590 [TROW])	True	1.294825	23.959022	-0.003987	False	4.0	4.
2018-01-02 00:00:00+00:00	Equity(23269 [WW])	True	1.481039	21.773177	-0.064942	False	6.0	6.
	Equity(24124 [WYNN])	True	1.280397	23.575596	0.000831	False	4.0	6.
	Equity(24811 [GES])	True	1.272976	21.051696	-0.014594	False	6.0	6.
	Equity(24832 [RL])	True	1.206513	22.854768	0.010624	False	4.0	4.
	Equity(27676 [AMP])	True	1.234680	23.944926	-0.012584	False	4.0	4.
	Equity(32902 [FSLR])	True	1.559013	22.676480	-0.012723	False	6.0	6.

In [48]: X_1 = pipeline_output_2[feature_cols]
y_1 = pipeline_output_2["return"]

```
In [49]: X_train_1,X_test_1,y_train_1,y_test_1=train_test_split(X_1,y_1,test_size=0.2,random_state= None)
```

Fit New Logistic Model

```
In [50]: # instantiate the model (using the default parameters)
logmodel_2 = LogisticRegression()
# fit the model with data
logmodel_2.fit(X_train_1,y_train_1)
```

```
Out[50]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
        intercept_scaling=1, max_iter=100, multi_class='ovr',
        penalty='l2', random_state=None, solver='liblinear', tol=0.0001,
        verbose=0)
```

```
In [51]: y_pred_1=logmodel.predict(X_test_1)
y_pred_1
```

```
Out[51]: array([ 1., -1.,  1., ..., -1.,  1., -1.])
```

```
In [52]: cnf_matrix_1 = metrics.confusion_matrix(y_test_1, y_pred_1)
print(cnf_matrix_1)
print("Accuracy:",metrics.accuracy_score(y_test_1, y_pred_1))
print("Precision:",metrics.precision_score(y_test_1, y_pred_1))
print("Recall:",metrics.recall_score(y_test_1, y_pred_1))
print(classification_report(y_test_1,y_pred_1))
```

```
[[235 458]
 [197 504]]
Accuracy: 0.530129124821
Precision: 0.523908523909
Recall: 0.718972895863
```

	precision	recall	f1-score	support
-1.0	0.54	0.34	0.42	693
1.0	0.52	0.72	0.61	701
avg / total	0.53	0.53	0.51	1394

Fit New Random Forest Classifier Model

```
In [53]: rfc_2 = RandomForestClassifier(n_estimators=100, max_leaf_nodes=2,max_features
      =7, min_samples_split=1)
      rfc_2.fit(X_train_1, y_train_1)
```

```
Out[53]: RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
      max_depth=None, max_features=7, max_leaf_nodes=2,
      min_samples_leaf=1, min_samples_split=1,
      min_weight_fraction_leaf=0.0, n_estimators=100, n_jobs=1,
      oob_score=False, random_state=None, verbose=0,
      warm_start=False)
```

```
In [54]: rfc_pred_2 = rfc.predict(X_test_1)
      rfc_pred_2
      cnf_matrix_1 = metrics.confusion_matrix(y_test_1, rfc_pred_2)
      print cnf_matrix_1
      print "Accuracy:",metrics.accuracy_score(y_test_1, rfc_pred_2)
      print "Precision:",metrics.precision_score(y_test_1, rfc_pred_2)
      print "Recall:",metrics.recall_score(y_test_1, rfc_pred_2)
      print classification_report(y_test_1,rfc_pred_2)
```

```
[[246 447]
 [212 489]]
Accuracy: 0.527259684362
Precision: 0.522435897436
Recall: 0.69757489301
```

	precision	recall	f1-score	support
-1.0	0.54	0.35	0.43	693
1.0	0.52	0.70	0.60	701
avg / total	0.53	0.53	0.51	1394

Out of Sample Predictions:

```
In [55]: pipeline_outsample_2 = run_pipeline(
      make_pipeline(),
      start_date="2018-01-01",
      end_date="2019-01-01"
      )
```

```
In [56]: pipeline_outsample_2['return'] = pipeline_outsample_2["return"].apply(np.sign)
      pipeline_outsample_2['Market Capital.'] = pipeline_outsample_2["Market Capita
      l."].apply(np.log)

      pipeline_outsample_2["Trading Signal"] = pd.get_dummies(pipeline_outsample_2[
      'BUY'],drop_first=True)
```

In [57]: `pipeline_outsample_2.head(20)`

Out[57]:

		BUY	MAD	Market Capital.	Returns	SHORT	Sentdex	Sentdex_la
	Equity(114 [ADBE])	True	1.215385	25.178645	-0.001424	False	6.0	6.
	Equity(1267 [CAT])	True	1.331272	25.263894	-0.005301	False	6.0	6.
	Equity(1539 [CI])	True	1.202831	24.636723	-0.010037	False	5.0	5.
	Equity(1941 [CTAS])	True	1.204947	23.532140	-0.004980	False	4.0	4.
	Equity(3321 [GPS])	True	1.343680	23.306846	-0.010456	False	5.0	5.
	Equity(6546 [ROST])	True	1.221500	24.146460	-0.004094	False	6.0	6.
2018-01-02 00:00:00+00:00	Equity(7590 [TROW])	True	1.294825	23.959022	-0.003987	False	4.0	4.
	Equity(23269 [WW])	True	1.481039	21.773177	-0.064942	False	6.0	6.
	Equity(24124 [WYNN])	True	1.280397	23.575596	0.000831	False	4.0	6.
	Equity(24811 [GES])	True	1.272976	21.051696	-0.014594	False	6.0	6.
	Equity(24832 [RL])	True	1.206513	22.854768	0.010624	False	4.0	4.
	Equity(27676 [AMP])	True	1.234680	23.944926	-0.012584	False	4.0	4.
	Equity(32902 [FSLR])	True	1.559013	22.676480	-0.012723	False	6.0	6.
	Equity(114 [ADBE])	True	1.211652	25.192585	0.013634	False	6.0	6.
	Equity(1267 [CAT])	True	1.334898	25.260461	-0.003743	False	6.0	6.
	Equity(1941 [CTAS])	True	1.203180	23.538664	0.006417	False	4.0	4.
2018-01-03 00:00:00+00:00	Equity(2127 [DE])	True	1.263072	24.654567	0.009070	False	6.0	6.
	Equity(2298 [DHI])	True	1.377207	23.678109	-0.000587	False	6.0	6.
	Equity(3321 [GPS])	True	1.344450	23.301252	0.000591	False	5.0	5.
	Equity(5121 [MU])	True	1.326128	24.645165	0.061512	False	6.0	6.

In [58]: `X_ofs18 = pipeline_outsample_2[feature_cols]`
`y_ofs18 = pipeline_outsample_2["return"]`


```
In [59]: y_pred_ofs18 = logmodel_2.predict(X_ofs18)
y_pred_ofs18
```

```
Out[59]: array([ 1.,  1.,  1., ...,  1.,  1., -1.])
```

Logestic Regression

```
In [60]: cnf_matrix_ofs2 = metrics.confusion_matrix(y_ofs18, y_pred_ofs18)
print cnf_matrix_ofs2
print "Accuracy:",metrics.accuracy_score(y_ofs18, y_pred_ofs18)
print "Precision:",metrics.precision_score(y_ofs18, y_pred_ofs18)
print "Recall:",metrics.recall_score(y_ofs18, y_pred_ofs18)
print classification_report(y_ofs18,y_pred_ofs18)
```

```
[[ 174 1479]
 [ 129 1841]]
Accuracy: 0.556168920784
Precision: 0.554518072289
Recall: 0.934517766497
```

	precision	recall	f1-score	support
-1.0	0.57	0.11	0.18	1653
1.0	0.55	0.93	0.70	1970
avg / total	0.56	0.56	0.46	3623

Random Forest Classifier

```
In [61]: y_rfc_pred_ofs18 = rfc_2.predict(X_ofs18)
cnf_matrix_ofs2 = metrics.confusion_matrix(y_ofs18, y_rfc_pred_ofs18)
print cnf_matrix_ofs2
print "Accuracy:",metrics.accuracy_score(y_ofs18, y_rfc_pred_ofs18)
print "Precision:",metrics.precision_score(y_ofs18, y_rfc_pred_ofs18)
print "Recall:",metrics.recall_score(y_ofs18, y_rfc_pred_ofs18)
print classification_report(y_ofs18,y_rfc_pred_ofs18)
```

```
[[ 114 1539]
 [  91 1879]]
Accuracy: 0.550096605023
Precision: 0.549736688122
Recall: 0.953807106599
```

	precision	recall	f1-score	support
-1.0	0.56	0.07	0.12	1653
1.0	0.55	0.95	0.70	1970
avg / total	0.55	0.55	0.44	3623

```
In [62]: print('Number of securities that passed the filter: %d' % len(pipeline_output.  
index.levels[1].unique()))  
print('Number of securities that passed the filter: %d' % len(pipeline_output_  
2.index.levels[1].unique()))  
print('Number of securities that passed the filter: %d' % len(pipeline_outsamp  
le.index.levels[1].unique()))  
print('Number of securities that passed the filter: %d' % len(pipeline_outsamp  
le_2.index.levels[1].unique()))
```

```
Number of securities that passed the filter: 365  
Number of securities that passed the filter: 325  
Number of securities that passed the filter: 129  
Number of securities that passed the filter: 162
```

Use Models to Predict EOD Return's Direction:

Date: 05/06/2019

Wait till 7am ET for sentiment datasets to be updated

PsychSignal Trader Mood : Update Frequency: Daily (updated every morning at ~7am ET)

Sentdex Sentiment : Update Frequency: Daily (updated every morning at ~7am ET)

US Equities Pricing : Update Frequency: Daily (updated overnight after each trading day).

```
In [18]: # Prices Update Frequency: Daily (updated overnight after each trading day).  
  
predict = run_pipeline(  
    make_pipeline(),  
    start_date="2019-07-22",  
    end_date="2019-07-22"  
)
```

Pipeline Execution Time: 2 Minutes, 11.41 Seconds

In [19]: predict

Out[19]:

		BUY	MAD	Market Capital.	Returns	SHORT	Sentdex	Sentde
2019-07-22 00:00:00+00:00	Equity(460 [APD])	True	1.279344	4.942636e+10	-0.004389	False	4.0	
	Equity(4488 [LM])	True	1.236497	3.298690e+09	-0.006261	False	6.0	
	Equity(6683 [SBUX])	True	1.296297	1.093714e+11	-0.013213	False	3.0	
	Equity(7684 [TSN])	True	1.237770	2.947701e+10	-0.001976	False	4.0	
	Equity(7998 [VMC])	True	1.202992	1.796585e+10	-0.003005	False	6.0	
	Equity(16511 [KMX])	True	1.213953	1.404833e+10	-0.020663	False	6.0	
	Equity(23438 [GME])	False	0.452238	4.418018e+08	0.030952	True	-1.0	
	Equity(32902 [FSLR])	True	1.247712	7.009147e+09	0.008489	False	6.0	
	Equity(34395 [LULU])	True	1.221689	2.448399e+10	-0.007551	False	5.0	
	Equity(38936 [DG])	True	1.207936	3.649585e+10	-0.010504	False	5.0	

```
In [20]: predict['Market Capital.']= predict["Market Capital."].apply(np.log)

predict["Trading Signal"] = pd.get_dummies(predict['BUY'],drop_first=True)

predict
```

Out[20]:

		BUY	MAD	Market Capital.	Returns	SHORT	Sentdex	Sentdex_k
2019-07-22 00:00:00+00:00	Equity(460 [APD])	True	1.279344	24.623750	-0.004389	False	4.0	4
	Equity(4488 [LM])	True	1.236497	21.916791	-0.006261	False	6.0	6
	Equity(6683 [SBUX])	True	1.296297	25.418015	-0.013213	False	3.0	6
	Equity(7684 [TSN])	True	1.237770	24.106876	-0.001976	False	4.0	4
	Equity(7998 [VMC])	True	1.202992	23.611739	-0.003005	False	6.0	6
	Equity(16511 [KMX])	True	1.213953	23.365769	-0.020663	False	6.0	4
	Equity(23438 [GME])	False	0.452238	19.906372	0.030952	True	-1.0	-3
	Equity(32902 [FSLR])	True	1.247712	22.670482	0.008489	False	6.0	5
	Equity(34395 [LULU])	True	1.221689	23.921285	-0.007551	False	5.0	5
	Equity(38936 [DG])	True	1.207936	24.320464	-0.010504	False	5.0	5

In []:

```
In [21]: X_live = predict[feature_cols]
print logmodel.predict(X_live)
# print logmodel_2.predict(X_live)
print rfc.predict(X_live)
# print rfc_2.predict(X_live)
print logmodel.predict_proba(X_live)
# print logmodel_2.predict_proba(X_live)
print rfc.predict_proba(X_live)
# print rfc_2.predict_proba(X_live)
print sv_model.predict(X_live)
```

```
[ 1.  1.  1.  1.  1.  1. -1.  1.  1.  1.]
[ 1.  1.  1.  1.  1.  1. -1.  1.  1.  1.]
[[ 0.4393584  0.5606416 ]
 [ 0.46324522 0.53675478]
 [ 0.42302052 0.57697948]
 [ 0.43393543 0.56606457]
 [ 0.45124281 0.54875719]
 [ 0.45998763 0.54001237]
 [ 0.61874917 0.38125083]
 [ 0.46629742 0.53370258]
 [ 0.45695699 0.54304301]
 [ 0.4397027  0.5602973  ]]
[[ 0.43975568 0.56024432]
 [ 0.43975568 0.56024432]
 [ 0.46369274 0.53630726]
 [ 0.43975568 0.56024432]
 [ 0.43975568 0.56024432]
 [ 0.43975568 0.56024432]
 [ 0.58097924 0.41902076]
 [ 0.43975568 0.56024432]
 [ 0.44003627 0.55996373]
 [ 0.43975568 0.56024432]]
[ 1.  1.  1.  1.  1.  1. -1.  1.  1.  1.]
```

```
In [22]: predict.index.levels[1]
```

```
Out[22]: Index([ Equity(460 [APD]), Equity(4488 [LM]), Equity(6683 [SBUX]),
               Equity(7684 [TSN]), Equity(7998 [VMC]), Equity(16511 [KMX]),
               Equity(23438 [GME]), Equity(32902 [FSLR]), Equity(34395 [LULU]),
               Equity(38936 [DG])],
              dtype='object')
```

```
In [23]: predict.index.levels[1]
predict['Return Predictions LR'] = logmodel.predict(X_live)
predict["Return Predictions RFC"] = rfc.predict(X_live)
predict["Return Predictions SVM"] = sv_model.predict(X_live)
```

```
In [24]: drop_cols = ["BUY", "Returns", "SHORT", "close_price", "return"]
predict.drop(feature_cols, axis = 1, inplace=True)
predict.drop(drop_cols, axis = 1, inplace = True)
```

```
In [25]: predict["LR Probability (-1)"] = logmodel.predict_proba(X_live)[: ,0]
predict["LR Probability (1)"] = logmodel.predict_proba(X_live)[: ,1]
predict["RFC Probability (-1)"] = rfc.predict_proba(X_live)[: ,0]
predict["RFC Probability (1)"] = rfc.predict_proba(X_live)[: ,1]
```

```
In [26]: predict
```

Out[26]:

		Return Predictions LR	Return Predictions RFC	Return Predictions SVM	LR Probability (-1)	LR Probability (1)	Prob:
2019-07-22 00:00:00+00:00	Equity(460 [APD])	1.0	1.0	1.0	0.439358	0.560642	0.4
	Equity(4488 [LM])	1.0	1.0	1.0	0.463245	0.536755	0.4
	Equity(6683 [SBUX])	1.0	1.0	1.0	0.423021	0.576979	0.4
	Equity(7684 [TSN])	1.0	1.0	1.0	0.433935	0.566065	0.4
	Equity(7998 [VMC])	1.0	1.0	1.0	0.451243	0.548757	0.4
	Equity(16511 [KMX])	1.0	1.0	1.0	0.459988	0.540012	0.4
	Equity(23438 [GME])	-1.0	-1.0	-1.0	0.618749	0.381251	0.5
	Equity(32902 [FSLR])	1.0	1.0	1.0	0.466297	0.533703	0.4
	Equity(34395 [LULU])	1.0	1.0	1.0	0.456957	0.543043	0.4
	Equity(38936 [DG])	1.0	1.0	1.0	0.439703	0.560297	0.4

• NOTES: :

- Causality Test between stock closing prices and lagged sentdex sentiment score. Check how many lags to apply.
- 1: <https://github.com/statsmodels/statsmodels/blob/master/statsmodels/tsa/stattools.py>
(<https://github.com/statsmodels/statsmodels/blob/master/statsmodels/tsa/stattools.py>)
- 2: from statsmodels.tsa.stattools import grangercausalitytests

```

In [70]: def grangercausalitytests(x, maxlag, addconst=True, verbose=True):
    """four tests for granger non causality of 2 timeseries
    all four tests give similar results
    `params_ftest` and `ssr_ftest` are equivalent based on F test which is
    identical to lmtest:grangertest in R
    Parameters
    -----
    x : array, 2d
        data for test whether the time series in the second column Granger
        causes the time series in the first column
    maxlag : integer
        the Granger causality test results are calculated for all lags up to
        maxlag
    verbose : bool
        print results if true
    Returns
    -----
    results : dictionary
        all test results, dictionary keys are the number of lags. For each
        lag the values are a tuple, with the first element a dictionary with
        teststatistic, pvalues, degrees of freedom, the second element are
        the OLS estimation results for the restricted model, the unrestricted
        model and the restriction (contrast) matrix for the parameter f_test.
    Notes
    -----
    TODO: convert to class and attach results properly
    The Null hypothesis for grangercausalitytests is that the time series in
    the second column, x2, does NOT Granger cause the time series in the first
    column, x1. Grange causality means that past values of x2 have a
    statistically significant effect on the current value of x1, taking past
    values of x1 into account as regressors. We reject the null hypothesis
    that x2 does not Granger cause x1 if the pvalues are below a desired size
    of the test.
    The null hypothesis for all four test is that the coefficients
    corresponding to past values of the second time series are zero.
    'params_ftest', 'ssr_ftest' are based on F distribution
    'ssr_chi2test', 'lrtest' are based on chi-square distribution
    References
    -----
    http://en.wikipedia.org/wiki/Granger_causality
    Greene: Econometric Analysis
    """
    from scipy import stats

    x = np.asarray(x)

    if x.shape[0] <= 3 * maxlag + int(addconst):
        raise ValueError("Insufficient observations. Maximum allowable "
                           "lag is {0}".format(int((x.shape[0] - int(addconst))
                                                    /
                                                    (3) - 1)))

    resli = {}

    for mlg in range(1, maxlag + 1):
        result = {}

```

```

if verbose:
    print('\nGranger Causality')
    print('number of lags (no zero)', mlg)
mxlg = mlg

# create lagmat of both time series
dta = lagmat2ds(x, mxlg, trim='both', dropex=1)

#add constant
if addconst:
    dtaown = add_constant(dta[:, 1:(mxlg + 1)], prepend=False)
    dtajoint = add_constant(dta[:, 1:], prepend=False)
else:
    raise NotImplementedError('Not Implemented')
    #dtaown = dta[:, 1:mxlg]
    #dtajoint = dta[:, 1:]

# Run ols on both models without and with lags of second variable
res2down = OLS(dta[:, 0], dtaown).fit()
res2djoint = OLS(dta[:, 0], dtajoint).fit()

#print results
#for ssr based tests see:
#http://support.sas.com/rnd/app/examples/ets/granger/index.htm
#the other tests are made-up

# Granger Causality test using ssr (F statistic)
fgc1 = ((res2down.ssr - res2djoint.ssr) /
        res2djoint.ssr / mxlg * res2djoint.df_resid)
if verbose:
    print('ssr based F test:          F=%-8.4f, p=%-8.4f, df_denom=%d, '
          ' df_num=%d' % (fgc1,
                          stats.f.sf(fgc1, mxlg,
                                      res2djoint.df_resid),
                                      res2djoint.df_resid, mxlg))
result['ssr_ftest'] = (fgc1,
                      stats.f.sf(fgc1, mxlg, res2djoint.df_resid),
                      res2djoint.df_resid, mxlg)

# Granger Causality test using ssr (ch2 statistic)
fgc2 = res2down.nobs * (res2down.ssr - res2djoint.ssr) / res2djoint.ss

if verbose:
    print('ssr based chi2 test:  chi2=%-8.4f, p=%-8.4f, '
          'df=%d' % (fgc2, stats.chi2.sf(fgc2, mxlg), mxlg))
result['ssr_chi2test'] = (fgc2, stats.chi2.sf(fgc2, mxlg), mxlg)

#Likelihood ratio test pvalue:
lr = -2 * (res2down.llf - res2djoint.llf)
if verbose:
    print('likelihood ratio test: chi2=%-8.4f, p=%-8.4f, df=%d' %
          (lr, stats.chi2.sf(lr, mxlg), mxlg))
result['lrtest'] = (lr, stats.chi2.sf(lr, mxlg), mxlg)

# F test that all lag coefficients of exog are zero
rconstr = np.column_stack((np.zeros((mxlg, mxlg)),
                           np.eye(mxlg, mxlg),

```



```

        np.zeros((mxlg, 1))))
ftres = res2djoint.f_test(rconstr)
if verbose:
    print('parameter F test:          F=%-8.4f, p=%-8.4f, df_denom=%d, '
          ' df_num=%d' % (ftres.fvalue, ftres.pvalue, ftres.df_denom,
                          ftres.df_num))
result['params_ftest'] = (np.squeeze(ftres.fvalue)[()],
                          np.squeeze(ftres.pvalue)[()],
                          ftres.df_denom, ftres.df_num)

resli[mxlg] = (result, [res2down, res2djoint, rconstr])

return resli

```

In []:

In []:

In []:

In []:

```

In [130]: beg_date = '2019-06-01'
          end_date = '2019-07-01'

```

```


In [131]: stock = get_pricing("NDAQ", start_date = beg_date,
                              end_date = end_date,
                              frequency = 'daily')

```

```

In [132]: stock['close_price'].plot(label = "Closing Price", figsize = (20,8), c = 'blue',
                                     marker='o',
                                     markerfacecolor='red', markersize=6, linestyle = "--")
plt.xlabel("Date")
plt.ylabel("Closing Price")
plt.title("AMAZON Clsong Prices 2018-07-01 till 2019-07-01)")
plt.legend(loc = 2);

```



```
In [133]: from quantopian.interactive.data.sentdex import sentiment as dataset

# import data operations
from odo import odo
```

```
In [134]: dataset.dshape
```

```
Out[134]: dshape("""var * {
    symbol: string,
    sentiment_signal: float64,
    sid: int64,
    asof_date: datetime,
    timestamp: datetime
}""")
```

```
In [135]: ndaq = dataset[dataset.symbol == "NDAQ"]
ndaq_df = odo(ndaq.sort('asof_date'), pd.DataFrame)
plt.plot(ndaq_df.asof_date, ndaq_df.sentiment_signal, marker='.', linestyle='None', color='r')
# plt.plot(ndaq_df.asof_date, pd.rolling_mean(ndaq_df.sentiment_signal, 21))
# plt.plot(ndaq_df.asof_date, pd.rolling_mean(ndaq_df.sentiment_signal, 252))
plt.xlabel("As Of Date (asof_date)")
plt.ylabel("Sentiment")
plt.title("Sentdex Sentiment for NDAQ")
plt.legend(["Sentiment - Single Day"], loc=1)
x1,x2,y1,y2 = plt.axis()
plt.axis((x1,x2,-4,7.5));
```



```
In [136]: initial = ndaq_df.index[ndaq_df.asof_date == beg_date][0]
end = ndaq_df.index[ndaq_df.asof_date == end_date][0] + 1
```

```
In [137]: my_test_df = pd.DataFrame()
```

```
In [138]: my_test_df["ClosePrice"] = stock['close_price']
```

```
In [139]: my_test_df["SentimentScore"] = ndaq_df['sentiment_signal'][initial:end]
```

```
In [143]: len(stock['close_price'])
```

```
Out[143]: 21
```

```
In [144]: ndaq_df.head()
```

```
Out[144]:
```

	symbol	sentiment_signal	sid	asof_date	timestamp
0	NDAQ	6.0	27026	2013-01-16	2013-01-17
1	NDAQ	6.0	27026	2013-01-17	2013-01-18
2	NDAQ	6.0	27026	2013-01-18	2013-01-19
3	NDAQ	6.0	27026	2013-01-19	2013-01-20
4	NDAQ	6.0	27026	2013-01-20	2013-01-21

```
In [147]: stock.index
```

```
Out[147]: DatetimeIndex(['2019-06-03', '2019-06-04', '2019-06-05', '2019-06-06',  
                        '2019-06-07', '2019-06-10', '2019-06-11', '2019-06-12',  
                        '2019-06-13', '2019-06-14', '2019-06-17', '2019-06-18',  
                        '2019-06-19', '2019-06-20', '2019-06-21', '2019-06-24',  
                        '2019-06-25', '2019-06-26', '2019-06-27', '2019-06-28',  
                        '2019-07-01'],  
                        dtype='datetime64[ns, UTC]', freq='C')
```

```
In [ ]:
```